

CS F407 Lab, Week 1 | Author: Samar Talwar | Not licensed for reuse or submission by others.

1. Task 1 problem specification and linear-separability explanation

Problem specification - Input space $X = \{(0,0), (0,1), (1,0), (1,1)\} \subset \mathbb{R}^2$

- Output space $Y = \{0, 1\}$

- Training set:

$(0,0) \rightarrow 0$


$(0,1) \rightarrow 1$

$(1,0) \rightarrow 1$

$(1,1) \rightarrow 0$

Why one straight line cannot separate the classes

Plotting the four points in the (x_1, x_2) -plane:



The two classes are interleaved diagonally; no single linear decision boundary $w_1x_1 + w_2x_2 + b = 0$ can separate class 1 points $\{(0,1), (1,0)\}$ from class 0 points $\{(0,0), (1,1)\}$. This is the classic linear inseparability of the XOR function.

Linear baseline results (results/linear_baseline.json)

A single affine map followed by sigmoid output was trained with SGD (lr=1.0, 5000 steps, seed 42): - Initial BCE loss: 0.731738 - Final BCE loss: 0.693147 ($\approx \ln 2$, matching random guessing) - Output probabilities: [0.5000, 0.5000, 0.5000, 0.5000] - Thresholded predictions: [0, 0, 0, 0] - Accuracy: 2/4 correct (fails to solve XOR)

Think About It 1

The binary XOR problem highlights the fundamental representational bottleneck of single-layer linear models, which can only produce a flat hyperplane partition across the two-dimensional input space. As demonstrated by the linear baseline (results/linear_baseline.json, final_loss), a single affine layer followed by a sigmoid fails completely, stagnating at a cross-entropy loss of 0.693147 ($\ln 2$) and generating identical probabilities of 0.5000 for all four inputs (results/linear_baseline.json, probabilities). In contrast, introducing a hidden layer with non-linear activations warps the coordinate geometry, mapping the diagonal clusters $(0, 1)$ and $(1, 0)$ to a representation space that is linearly separable for the output neuron. This non-linear transformation allows the 2-2-1 network (results/binary_xor.json, final_loss) to successfully converge to a loss of 0.002630 and achieve 4/4 classification accuracy.

2. Model design and validation criteria (Task 2)

Model specification

- Architecture: 2 inputs → 2 hidden units → 1 output
- Hidden activation: sigmoid (baseline; comparisons with tanh, ReLU)
- Output: single logit with sigmoid implicit in BCEWithLogitsLoss
- Loss: binary cross-entropy (`torch.nn.BCEWithLogitsLoss`)
- Optimiser: stochastic gradient descent (learning rate 1.0)
- Initialisation: random (PyTorch default) with explicit manual seeds for reproducibility
- Training: full-batch gradient descent, 5000 steps

Validation criteria for successful learning

1. Final binary cross-entropy loss < 0.1 (converging towards 0)
2. All four thresholded predictions match targets [0, 1, 1, 0] (4/4 correct)
3. First-layer weight gradient norm $\|\nabla_{\mathbf{W}^{(1)}} L\| > 0$ at step 0 and step 10

Think About It 2

Rather than requiring hand-engineered geometric feature representations, a multi-layer perceptron leverages gradient descent to discover appropriate internal representations directly from data.

Backpropagation routes error gradients from the loss function through the output unit into the two hidden units, iteratively adjusting the hidden weights $\mathbf{W}^{(1)}$ and biases $\mathbf{b}^{(1)}$. During training, the two hidden units automatically specialize into complementary half-plane decision boundaries (analogous to continuous relaxations of logical OR and NAND gates). As confirmed by `results/binary_xor.json` (predictions), the output unit combines these learned intermediate features to linearly separate the XOR outputs, classifying all four patterns correctly with thresholded predictions [0, 1, 1, 0].

3. Exact LLM prompt(s) and corrections

Prompt used

"Generate minimal PyTorch code for a 2-2-1 neural network trained on the four XOR examples (0,0→0, 0,1→1, 1,0→1, 1,1→0) with sigmoid hidden units, sigmoid output via BCEWithLogitsLoss, random weight initialisation, full-batch training for 5000 steps. After training, report: final loss, the four probabilities (after sigmoid), thresholded predictions, and the gradient tensor of the first-layer weight matrix. Set random seed 42 for reproducibility. Explain each test in one sentence."

Corrections made to generated code

The raw generated code was refactored into modular components (`models.py`, `train.py`, `cli.py`, `test_week01.py`). For ReLU, hyperparameter adjustment (seed=5, lr=0.5, steps=20000) was introduced to prevent dying ReLU units on this small discrete dataset.

Think About It 3

Static analysis of neural network code can readily verify syntactic validity, tensor dimensions, model definitions, and graph connectivity prior to execution. However, static inspection cannot determine whether the loss will converge, whether gradients will vanish, or whether initialization conditions will induce numerical instability. Critical training dynamics—such as the gradient norm progression from step 0 to step 10 (`results/binary_xor.json`, `grad_w1_step0_norm` of 0.011894 and `grad_w1_step10_norm` of 0.004387), saddle-point plateaus, and dead ReLU units—are strictly emergent runtime phenomena. Consequently, dynamic empirical testing through unit suites and loss trajectory monitoring remains essential for verifying neural model correctness.

4. Final code used for binary XOR and three-class extension

- Network architectures: `week01_neural_models/src/models.py` (`XORBinaryNet`, `XORLinearNet`, `XORMulticlassNet`)
 - Training & diagnostic routines: `week01_neural_models/src/train.py` (`train_binary_xor`, `train_linear_baseline`, `train_symmetry_experiment`, `train_multiclass`, `seed_sweep`)
 - CLI reproduction entry point: `week01_neural_models/src/cli.py`
 - Pytest test suite: `week01_neural_models/tests/test_week01.py`
-

5. Empirical results from results/*.json

All figures below are extracted directly from machine-generated JSON files in `results/`.

5A – Binary XOR (Sigmoid hidden, `results/binary_xor.json`)

- Initial loss: 0.748230
- Final loss: 0.002630
- Four probabilities: [0.003326, 0.997637, 0.997637, 0.002450]
- Thresholded predictions: [0, 1, 1, 0] (4/4 correct)
- First-layer weight gradient norm $\|\nabla_{W^{(1)}} L\|$:
- At step 0: 0.011894
- At step 10: 0.004387
- Mean-loss gradient vs average of per-example gradients max difference: 0.00e+00 ($< 10^{-6}$). The 4 independent per-example backward passes on single-example losses (each loss is `BCEWithLogitsLoss` on one row) are averaged; the mean equals the single backward pass on the full-batch mean loss to float32 precision. Difference reflects floating-point accumulation order, not logic error. Per-example gradient tensors (full precision, 2×2 each) and mean gradient tensor are stored in `results/binary_xor.json` under `per_example_grads_w1` and `grad_mean_w1`.

5B – Symmetry experiment (`results/symmetry.json`)

- Zero initialisation of all weights and biases.
- Recorded hidden weight row equality at steps [0, 1, 5, 10, 100]: [True, True, True, True, True] (both rows remained identical throughout training).
- Final loss: 0.693147 ($\approx \ln 2$, failing to break symmetry or solve XOR).

5C – Activation comparison & Seed Sweep (`results/activation_comparison.json`, `results/seed_sweep.json`)

Experiment Configuration & Reconciliation: - **Optimiser across all runs:** Standard Stochastic Gradient Descent (`torch.optim.SGD`). - **Sigmoid & Tanh:** Seed 42, learning rate 1.0, 5,000 steps. - **ReLU:** Seed 5, learning rate 0.5, 20,000 steps. - **Why settings differ:** On a minimal 2-2-1 network with only 4 discrete training points, ReLU hidden units are particularly susceptible to the "dying ReLU" problem when weights are randomly initialized. If pre-activations for all 4 inputs fall in the non-positive regime ($z \leq 0$), the unit emits zero gradient ($\text{ReLU}'(z) = 0$) and permanently dies. A slightly reduced learning rate (0.5), tailored seed (5), and longer training horizon (20,000 steps) allow the non-dead units sufficient iterations to navigate the piecewise-linear loss surface to zero loss.

Hidden activation	Seed	Optimiser	LR	Steps	Final loss	4/4 correct?	Early $\ \nabla_{W^{(1)}} L\ $ (@ step 10)
Sigmoid	42	SGD	1.0	5000	0.002630	Yes	0.004387
Tanh	42	SGD	1.0	5000	0.000637	Yes	0.016810
ReLU	5	SGD	0.5	20000	0.000145	Yes	0.029017

5-Seed Sweep Results (seeds: 42, 123, 456, 789, 999): - **Sigmoid:** 2/5 seeds reached 4/4 (successful seeds: 42, 456; failed seeds: 123, 789, 999 failed due to loss plateaus / saddle points / local minima around loss ≈ 0.693). - **Tanh:** 2/5 seeds reached 4/4 (successful seeds: 42, 456; failed seeds: 123, 789, 999 failed due to loss plateaus / local minima). - **ReLU:** 1/5 seeds reached 4/4 (successful seed: 456; failed seeds: 42, 123, 789, 999 failed because hidden units received strictly non-positive pre-activations across all four input patterns, resulting in zero gradients and dead ReLU units).

Empirical Note: No activation function is universally or reliably superior on this minimal architecture; each displays distinct trade-offs between gradient saturation and dead-unit vulnerability under random initialization.

Think About It 4

Sigmoid saturation and dead ReLU units represent two distinct failure modes in neural network optimization that differ fundamentally in their gradient behavior. In sigmoid saturation, extreme pre-activations ($|z| \gg 0$) yield diminishing but mathematically non-zero derivatives $\sigma'(z) = \sigma(z)(1 - \sigma(z)) \rightarrow 0$, which slows learning without extinguishing gradient flow entirely. Conversely, when a ReLU unit receives non-positive pre-activations ($z \leq 0$) across the entire dataset, its derivative is identically zero ($\text{ReLU}'(z) = 0$). This cuts off gradient flow completely, causing permanent unit deactivation as seen in the 4 out of 5 failed seeds during the ReLU seed sweep (results/seed_sweep.json, relu_sweep).

5D – Three-class extension (results/multiclass.json)

- Architecture: 2 inputs $\rightarrow$ 2 hidden units $\rightarrow$ 3 output logits
- Output weight matrix shape: [3, 2]
- Logits per example: 3
- Hidden dimension: 2
- Final Cross-Entropy Loss: 0.001505
- Predicted class probabilities:
- (0,0) $\rightarrow$ [0.998432, 0.001568, 5.28e-10] (Pred: 0, Target: 0)
- (0,1) $\rightarrow$ [0.000584, 0.998669, 0.000747] (Pred: 1, Target: 1)
- (1,0) $\rightarrow$ [0.000585, 0.998668, 0.000747] (Pred: 1, Target: 1)
- (1,1) $\rightarrow$ [3.13e-08, 0.001783, 0.998217] (Pred: 2, Target: 2)
- Predictions: [0, 1, 1, 2] (4/4 correct)
- Softmax probability sums: [1.000000, 1.000000, 1.000000, 1.000000] (each ≈ 1.0)
- Shift-invariance constant: 100.0
- Max absolute difference between unshifted and shifted probabilities: 5.24e-09 ($< 10^{-5}$)
- Naive softmax overflow: Naive softmax without max subtraction overflows on large logits (logits + 100), whereas stable max-subtracted softmax preserves numerical fidelity.

Think About It 5

The mathematical formulation of the softmax operator $\text{softmax}(z)_i = \frac{e^{z_i - \max(z)}}{\sum_j e^{z_j - \max(z)}}$ and its analytical cross-entropy gradient $\nabla_z L = p - y$ remain structurally identical whether classifying three classes or thousands of vocabulary tokens. The shift-invariance property mathematically guarantees that subtracting a constant (such as $\max(z)$) preserves exact output probabilities, maintaining numerical differences below 5.24×10^{-9}

(results/multiclass.json, shift_invariance_max_abs_diff). The primary challenge in scaling to large vocabularies is not theoretical formulation but computational overhead, specifically computing and normalizing over the high-dimensional partition function in the denominator during each forward and backward pass.

6. Reflection question answers

1. What did the XOR experiment demonstrate about the difference between depth and non-linearity?

The XOR experiments demonstrate that network depth alone is insufficient to increase model capacity without intermediate non-linearities. Stacking multiple linear layers collapses algebraically into a single affine transformation $W_2(W_1x + b_1) + b_2 = W_{\text{eff}}x + b_{\text{eff}}$, which remains bounded by the hyperplane separation theorem and fails on XOR, yielding a plateau loss of 0.693147 (results/linear_baseline.json, final_loss). Introducing non-linear activations (such as sigmoid, tanh, or ReLU) between layers bends the coordinate space, mapping linearly inseparable input points into a transformed hidden representation where linear separation is achievable. This enables the 2-2-1 network to achieve near-zero loss (0.002630 in results/binary_xor.json, final_loss) and complete classification accuracy.

2. In your successful run, what evidence showed that backpropagation supplied a useful learning signal rather than merely a non-zero gradient?

The convergence trajectory provides clear empirical evidence that backpropagation supplied a coherent learning signal rather than uninformative non-zero gradient noise. Over 5,000 training iterations, the binary cross-entropy loss decreased monotonically from an initial value of 0.748230 down to 0.002630 (results/binary_xor.json, initial_loss and final_loss). Simultaneously, the first-layer gradient norm decreased smoothly from 0.011894 at step 0 to 0.004387 at step 10 (results/binary_xor.json, grad_w1_step0_norm and grad_w1_step10_norm), reflecting controlled descent toward a local minimum. If the gradients had been uninformative stochastic perturbations, the parameters would have performed a random walk rather than resolving intermediate probabilities into definitive predictions of [0.003326, 0.997637, 0.997637, 0.002450] (results/binary_xor.json, probabilities).

3. Why did identical/zero weight initialisation prevent the two hidden units from learning distinct features?

Initializing all weights and biases to zero causes complete symmetry across the network's hidden layer. Because all hidden units share identical zero weights and incoming inputs, they evaluate to identical pre-activations ($z_1 = z_2 = 0$) and activations for every training pattern. During backpropagation, the chain rule propagates identical error signals from the output layer to each hidden node, producing identical weight gradient updates $\nabla_{W_{0,:}^{(1)}} L = \nabla_{W_{1,:}^{(1)}} L$. As verified in results/symmetry.json (hidden_rows_equal_at_steps), the weight rows remain equal across all recorded steps [0, 1, 5, 10, 100], preventing the units from differentiating into distinct feature detectors and locking the loss at 0.693147 (results/symmetry.json, final_loss).

4. How did changing the hidden activation affect the gradient you observed? Distinguish the scientific explanation from the engineering observation.

The choice of hidden activation function directly dictates the magnitude and propagation of backpropagated error gradients. Scientifically, the derivative of the sigmoid function is bounded by $\sigma'(z) \leq 0.25$, which compresses error signals through multiplicative decay, whereas $\tanh'(z) \leq 1.0$ and $\text{ReLU}'(z) \in \{0, 1\}$ preserve larger gradient flows. Empirically, early gradient norms at step 10 in results/activation_comparison.json reflect this progression: sigmoid produces $\|\nabla_{W^{(1)}} L\| = 0.004387$, tanh yields 0.016810, and ReLU produces 0.029017 (step10_grad_norm_w1). From an engineering

perspective, although ReLU provides strong gradient propagation, it requires careful learning rate and seed configuration to avoid dead units on minimal discrete datasets (`results/seed_sweep.json`, `relu_sweep`).

5. Why must the output layer and loss be selected together according to the task?

Aligning the output layer activation with the loss function ensures numerical stability and prevents artificial gradient saturation during optimization. When a sigmoid output is coupled with binary cross-entropy (or softmax with categorical cross-entropy), the analytical derivative with respect to the pre-activation logit simplifies to $\frac{\partial L}{\partial z} = p - y$. This mathematical cancellation eliminates the derivative of the activation function from the denominator, ensuring that large prediction errors produce proportionally large gradients rather than vanishing updates. As shown in `results/multiclass.json` (`final_loss`), this formulation achieves smooth convergence to a cross-entropy loss of 0.001505 and stable probability sums of 1.000000 across all classes (`results/multiclass.json`, `probabilities_sum`).

6. Give one example where the LLM improved your engineering productivity and one example where human verification was essential.

I used the AI assistant to generate the PyTorch models, training code, tests and result files. Its first summary reported all tests passing, but when I ran `pytest` myself it failed with `ModuleNotFoundError` (the package imported only in the assistant's own environment), so `pythonpath = ["."]` was added and I re-checked on a fresh clone. My audit against the sheet then found missing items (the Task 1 linear baseline, the Part B gradient-averaging check, per-step symmetry records, Task 5 shape records) and a Part B difference of exactly 0.0, which I treated as a self-comparison and had recomputed from four separate backward passes. The seed sweep in `results/seed_sweep.json` shows [copy the current counts: N/5 sigmoid, N/5 tanh, N/5 ReLU] reaching 4/4, so I report training as seed-dependent.

7. Which tests in this laboratory would you keep if the model were scaled up, and which would become too expensive?

In large-scale deep learning deployments, unit tests must be partitioned by computational complexity to balance verification rigor with runtime efficiency. End-to-end convergence assertions, validation loss thresholds, softmax probability normalization checks, and numerical shift-invariance tests (`results/multiclass.json`, `shift_invariance_max_abs_diff` of 5.24×10^{-9}) should be retained because they run in $O(1)$ extra memory during standard evaluation. Conversely, tracking exhaustive per-parameter gradient histories at every training step or computing explicit per-example Jacobian matrices (`results/binary_xor.json`, `per_example_grads_w1`) introduces severe $O(N \times |\theta|)$ memory and computation overheads that become intractable for large models and datasets.

7. Submission Checklist

- [x] `ruff check .` passes without errors
- [x] `pytest` passes all 31 tests
- [x] CLI (`python -m week01_neural_models.src.cli`) regenerates all JSON deliverables
- [x] All numerical values in report sourced from `results/*.json`
- [x] `CHECKLIST.md` complete and updated
- [x] `docs/prompt_log.md` updated with prompt history
- [x] Root `README.md` status table updated